

HOMework 1

MULTILAYER PERCEPTRONS AND ARIMA MODEL FOR TIME SERIES FORECASTING

BMEN 4470 - Deep Learning for Biomedical Signal Processing

Fall 2021

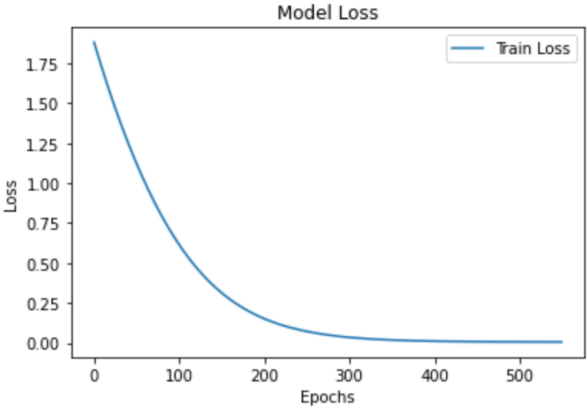
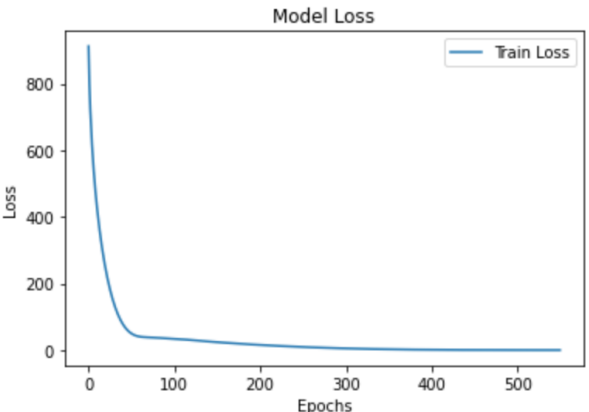
Michelle Campoli

mec2308

Problem 1a.

Accuracy of a perceptron is dependent on the weights selected and whether the model can be linearly or logistically fit. Increasing the number of epochs increases the accuracy of the prediction.

In the table below, a perceptron and a multi-layer perceptron are compared.

	Perceptron	Add 1 hidden layer															
Model	<p>Model: "sequential_23"</p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>dense_32 (Dense)</td><td>(None, 1)</td><td>4</td></tr> </table> <p> Total params: 4 Trainable params: 4 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	dense_32 (Dense)	(None, 1)	4	<p>Model: "sequential"</p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>dense (Dense)</td><td>(None, 8)</td><td>32</td></tr> <tr> <td>dense_1 (Dense)</td><td>(None, 1)</td><td>9</td></tr> </table> <p> Total params: 41 Trainable params: 41 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	dense (Dense)	(None, 8)	32	dense_1 (Dense)	(None, 1)	9
Layer (type)	Output Shape	Param #															
dense_32 (Dense)	(None, 1)	4															
Layer (type)	Output Shape	Param #															
dense (Dense)	(None, 8)	32															
dense_1 (Dense)	(None, 1)	9															
Loss																	
Prediction	<pre>#Demonstrate Prediction using model.predict x_input = np.array([70, 80, 90]) # predict the output x_input = x_input.reshape((1, n_steps)) y_output = model.predict(x_input) print("X=%s, Predicted=%s" % (x_input, y_output))</pre> <p>X=[[70 80 90]], Predicted=[[100.150566]]</p>	<pre>x_input = np.array([70, 80, 90]) # predict the out x_input = x_input.reshape((1, n_steps)) y_output = model.predict(x_input) print("X=%s, Predicted=%s" % (x_input, y_output))</pre> <p>X=[[70 80 90]], Predicted=[[101.54166]]</p>															

By adding hidden layers, the number of epochs required to minimize the loss decreases.

Problem 1b.

The model from 1a was used to predict the test sequence $[1^2, 2^2, 3^2]$.

```
x_input = np.array([1, 4, 9]) # predict the output
x_input = x_input.reshape((1, n_steps))

y_output = model.predict(x_input)
print("X=%s, Predicted=%s" % (x_input, y_output))

X=[[1 4 9]], Predicted=[[13.181818]]
```

The model from 1a. does not accurately predict the test sequence $[1^2, 2^2, 3^2] 4^2$, as shown above. The prediction is 13 and the actual value is 16. This is because the sequence is non-linear, and the model has been trained on data that is linearly separable. The model was also built for a linearly separable dataset as the weights implemented at each layer are the same ones used to optimize linear regression models.

Problem 2a.

Using Keras Functional API, a model was built to predict exponential growth. Three hidden layers were added to reduce the number of epochs necessary for an accurate prediction. Weights were also implemented at each layer. Given the non-linear nature of the dataset, the weights were similar to those used in logistic regression.

Model

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 3)]	0
dense (Dense)	(None, 256)	1024
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 1)	129
Total params: 50,561		
Trainable params: 50,561		
Non-trainable params: 0		

Loss

Prediction

```

x_input = np.array([49, 64, 81]) # predict the ooutput for this input
x_input = x_input.reshape((1, n_steps))
y_output = model.predict(x_input)
print("X=%s, Predicted=%s" % (x_input, y_output))

```

X=[[49 64 81]], Predicted=[[102.64433]]

The Functional API model was able to better predict the next exponential outcome with a 2% error (102.6 v. 100).

Problem 2b.

The model above has 3 inputs, 3 hidden layers with 256, 128, and 128 neurons, and 1 output. There are a total of 512 neurons. 4 weights were used to simulate a logistic regression, one weight per layer. There are 50,561 trainable parameters in this model.

Problem 3.

look_back	RMSE
7	Train RMSE = 1204 Test RMSE = 1774
14	Train RMSE = 1029 Test RMSE = 1893
30	Train RMSE = 898 Test RMSE = 1857
90	Train RMSE = 1333 Test RMSE = 2353

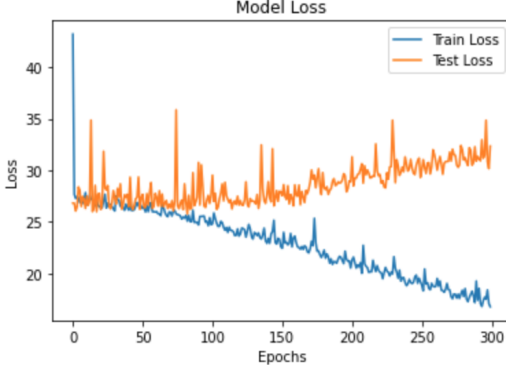
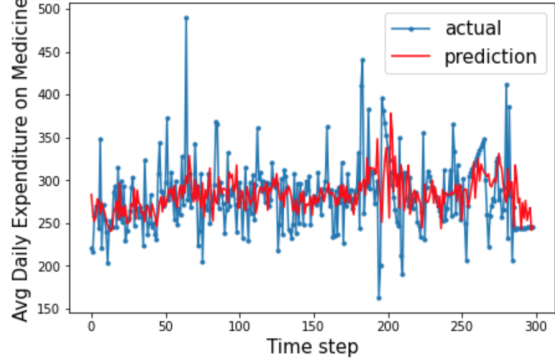
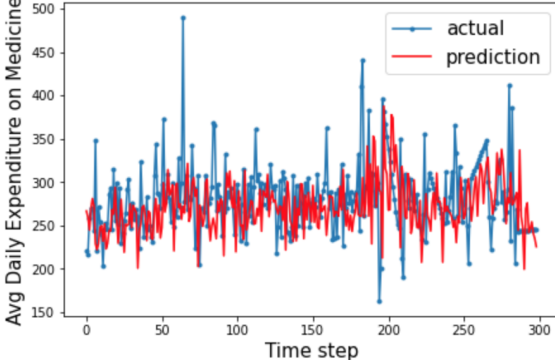
Different look-back time-steps were tested to achieve the lowest root mean square error (RMSE). The given look-back was 7 days, correlating to weekly medical expenditures. I also tested 14 days (biweekly), 30 days (monthly), and 90 days (quarterly) to see which was best to predict the trend in medical expenditures. I chose a model from Problem 2 and ran each look_back value listed. The RMSE for train and validation (test) datasets are shown. There was no clear best look-back time-step. Look_back = 14 was chosen because the train and test RMSEs were lowest without the model overfitting the data.

Experimental Models

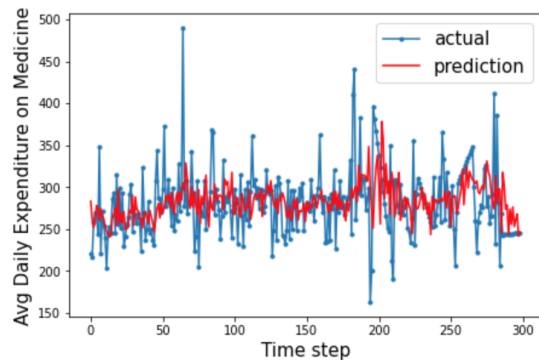
A model was designed that achieved the lowest RMSE given what was learned in Problems 1 and 2. The following three parameters will be changed:

1. Number of epochs
2. Number of hidden layers
3. Activation Function

to see their effects on the model's accuracy (RMSE) and loss. The results will be analyzed in the following three tables.

	Initial Model, Epoch = 150	Change Epoch = 300																																																						
Model	<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_2 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense_4 (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_5 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_6 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_7 (Dense)</td><td>(None, 1)</td><td>65</td></tr> <tr> <td colspan="3">Total params: 14,401</td></tr> <tr> <td colspan="3">Trainable params: 14,401</td></tr> <tr> <td colspan="3">Non-trainable params: 0</td></tr> </tbody> </table>	Layer (type)	Output Shape	Param #	input_2 (InputLayer)	[(None, 14)]	0	dense_4 (Dense)	(None, 128)	1920	dense_5 (Dense)	(None, 64)	8256	dense_6 (Dense)	(None, 64)	4160	dense_7 (Dense)	(None, 1)	65	Total params: 14,401			Trainable params: 14,401			Non-trainable params: 0			<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_2 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense_4 (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_5 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_6 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_7 (Dense)</td><td>(None, 1)</td><td>65</td></tr> <tr> <td colspan="3">Total params: 14,401</td></tr> <tr> <td colspan="3">Trainable params: 14,401</td></tr> <tr> <td colspan="3">Non-trainable params: 0</td></tr> </tbody> </table>	Layer (type)	Output Shape	Param #	input_2 (InputLayer)	[(None, 14)]	0	dense_4 (Dense)	(None, 128)	1920	dense_5 (Dense)	(None, 64)	8256	dense_6 (Dense)	(None, 64)	4160	dense_7 (Dense)	(None, 1)	65	Total params: 14,401			Trainable params: 14,401			Non-trainable params: 0		
Layer (type)	Output Shape	Param #																																																						
input_2 (InputLayer)	[(None, 14)]	0																																																						
dense_4 (Dense)	(None, 128)	1920																																																						
dense_5 (Dense)	(None, 64)	8256																																																						
dense_6 (Dense)	(None, 64)	4160																																																						
dense_7 (Dense)	(None, 1)	65																																																						
Total params: 14,401																																																								
Trainable params: 14,401																																																								
Non-trainable params: 0																																																								
Layer (type)	Output Shape	Param #																																																						
input_2 (InputLayer)	[(None, 14)]	0																																																						
dense_4 (Dense)	(None, 128)	1920																																																						
dense_5 (Dense)	(None, 64)	8256																																																						
dense_6 (Dense)	(None, 64)	4160																																																						
dense_7 (Dense)	(None, 1)	65																																																						
Total params: 14,401																																																								
Trainable params: 14,401																																																								
Non-trainable params: 0																																																								
Training	<pre>loss='mean_absolute_error' optimizer="adam" epochs=150, validation_split=0.2</pre>	<pre>loss='mean_absolute_error' optimizer="adam" epochs=300, validation_split=0.2</pre>																																																						
Loss																																																								
RMSE	Train RMSE: 23.317861557006836 Test RMSE: 32.44575119018555	Train RMSE: 19.65681266784668 Test RMSE: 37.33937454223633																																																						
Prediction v. Actual																																																								

Increasing epoch number decreases the RMSE and increases accuracy of predicted values. However, the model is overfitting, as seen by the test loss increasing relative to the train loss.

	Initial Model (3 Layers)	Add 2 hidden layers (5 Layers)																																										
Model	<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_2 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense_4 (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_5 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_6 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_7 (Dense)</td><td>(None, 1)</td><td>65</td></tr> </tbody> </table> <p> Total params: 14,401 Trainable params: 14,401 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	input_2 (InputLayer)	[(None, 14)]	0	dense_4 (Dense)	(None, 128)	1920	dense_5 (Dense)	(None, 64)	8256	dense_6 (Dense)	(None, 64)	4160	dense_7 (Dense)	(None, 1)	65	<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_1 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_1 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_2 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_3 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_4 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_5 (Dense)</td><td>(None, 1)</td><td>65</td></tr> </tbody> </table> <p> Total params: 22,721 Trainable params: 22,721 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	input_1 (InputLayer)	[(None, 14)]	0	dense (Dense)	(None, 128)	1920	dense_1 (Dense)	(None, 64)	8256	dense_2 (Dense)	(None, 64)	4160	dense_3 (Dense)	(None, 64)	4160	dense_4 (Dense)	(None, 64)	4160	dense_5 (Dense)	(None, 1)	65
Layer (type)	Output Shape	Param #																																										
input_2 (InputLayer)	[(None, 14)]	0																																										
dense_4 (Dense)	(None, 128)	1920																																										
dense_5 (Dense)	(None, 64)	8256																																										
dense_6 (Dense)	(None, 64)	4160																																										
dense_7 (Dense)	(None, 1)	65																																										
Layer (type)	Output Shape	Param #																																										
input_1 (InputLayer)	[(None, 14)]	0																																										
dense (Dense)	(None, 128)	1920																																										
dense_1 (Dense)	(None, 64)	8256																																										
dense_2 (Dense)	(None, 64)	4160																																										
dense_3 (Dense)	(None, 64)	4160																																										
dense_4 (Dense)	(None, 64)	4160																																										
dense_5 (Dense)	(None, 1)	65																																										
Training	<pre> loss='mean_absolute_error' optimizer="adam" epochs=150, validation_split=0.2 </pre>	<pre> loss='mean_absolute_error' optimizer="adam" epochs=150, validation_split=0.2 </pre>																																										
Loss																																												
RMSE	Train RMSE: 23.317861557006836 Test RMSE: 32.44575119018555	Train RMSE: 23.906476974487305 Test RMSE: 31.657373428344727																																										
Prediction v. Actual																																												

The RMSE does not change with increasing hidden layers. The train and test loss are similar, indicating that the model is neither over nor under fit. Adding hidden layers decreases the overfitting of the model, making it just right.

	Initial Model (Activation function= ReLU)	Change Activation Function to Swish																																				
Model	<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_2 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense_4 (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_5 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_6 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_7 (Dense)</td><td>(None, 1)</td><td>65</td></tr> </tbody> </table> <p> Total params: 14,401 Trainable params: 14,401 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	input_2 (InputLayer)	[(None, 14)]	0	dense_4 (Dense)	(None, 128)	1920	dense_5 (Dense)	(None, 64)	8256	dense_6 (Dense)	(None, 64)	4160	dense_7 (Dense)	(None, 1)	65	<table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_1 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense (Dense)</td><td>(None, 128)</td><td>1920</td></tr> <tr> <td>dense_1 (Dense)</td><td>(None, 64)</td><td>8256</td></tr> <tr> <td>dense_2 (Dense)</td><td>(None, 64)</td><td>4160</td></tr> <tr> <td>dense_3 (Dense)</td><td>(None, 1)</td><td>65</td></tr> </tbody> </table> <p> Total params: 14,401 Trainable params: 14,401 Non-trainable params: 0 </p>	Layer (type)	Output Shape	Param #	input_1 (InputLayer)	[(None, 14)]	0	dense (Dense)	(None, 128)	1920	dense_1 (Dense)	(None, 64)	8256	dense_2 (Dense)	(None, 64)	4160	dense_3 (Dense)	(None, 1)	65
Layer (type)	Output Shape	Param #																																				
input_2 (InputLayer)	[(None, 14)]	0																																				
dense_4 (Dense)	(None, 128)	1920																																				
dense_5 (Dense)	(None, 64)	8256																																				
dense_6 (Dense)	(None, 64)	4160																																				
dense_7 (Dense)	(None, 1)	65																																				
Layer (type)	Output Shape	Param #																																				
input_1 (InputLayer)	[(None, 14)]	0																																				
dense (Dense)	(None, 128)	1920																																				
dense_1 (Dense)	(None, 64)	8256																																				
dense_2 (Dense)	(None, 64)	4160																																				
dense_3 (Dense)	(None, 1)	65																																				
Training	<pre> loss='mean_absolute_error' optimizer="adam" epochs=150, validation_split=0.2 </pre>	<pre> loss='mean_absolute_error' optimizer="adam" epochs=150, validation_split=0.2 </pre>																																				
Loss																																						
RMSE	Train RMSE: 23.317861557006836 Test RMSE: 32.44575119018555	Train RMSE: 22.17965316772461 Test RMSE: 33.41248321533203																																				
Prediction v. Actual																																						

Activation function was originally set to “relu” for each hidden layer. For this experimental test they were all changed to “swish”. The loss function has less variance after the change, but there is a minimal difference in RMSE. It also appears that there is overfitting using swish, as the train and test loss seem to drift further apart.

Use Models from Problem 1 & 2:

The MLP models from Problems 1 and 2 were trained with the Medical Expenditures dataset. The following table shows their results.

	Model from Problem 1 (Sequential)	Model from Problem 2 (Functional API)																											
Model	<p>Model: "sequential_8"</p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>dense_154 (Dense)</td><td>(None, 8)</td><td>120</td></tr> <tr> <td>dense_155 (Dense)</td><td>(None, 1)</td><td>9</td></tr> </table> <p>Total params: 129 Trainable params: 129 Non-trainable params: 0</p>	Layer (type)	Output Shape	Param #	dense_154 (Dense)	(None, 8)	120	dense_155 (Dense)	(None, 1)	9	<p>Model: "model_38"</p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>input_39 (InputLayer)</td><td>[(None, 14)]</td><td>0</td></tr> <tr> <td>dense_156 (Dense)</td><td>(None, 256)</td><td>3840</td></tr> <tr> <td>dense_157 (Dense)</td><td>(None, 128)</td><td>32896</td></tr> <tr> <td>dense_158 (Dense)</td><td>(None, 128)</td><td>16512</td></tr> <tr> <td>dense_159 (Dense)</td><td>(None, 1)</td><td>129</td></tr> </table> <p>Total params: 53,377 Trainable params: 53,377 Non-trainable params: 0</p>	Layer (type)	Output Shape	Param #	input_39 (InputLayer)	[(None, 14)]	0	dense_156 (Dense)	(None, 256)	3840	dense_157 (Dense)	(None, 128)	32896	dense_158 (Dense)	(None, 128)	16512	dense_159 (Dense)	(None, 1)	129
Layer (type)	Output Shape	Param #																											
dense_154 (Dense)	(None, 8)	120																											
dense_155 (Dense)	(None, 1)	9																											
Layer (type)	Output Shape	Param #																											
input_39 (InputLayer)	[(None, 14)]	0																											
dense_156 (Dense)	(None, 256)	3840																											
dense_157 (Dense)	(None, 128)	32896																											
dense_158 (Dense)	(None, 128)	16512																											
dense_159 (Dense)	(None, 1)	129																											
Loss																													
RMSE	<p>Train RMSE: 1235.6380615234375 Test RMSE: 1710.6234130859375</p>	<p>Train RMSE: 1129.9097900390625 Test RMSE: 1859.9356689453125</p>																											
Prediction v. Actual																													

Models from problems 1 and 2 return high RMSE values whose loss does not improve with increasing epochs. The graphs also show that the predicted outputs do not align with the actual data. These are not good models for predicting the Medical Expenditures dataset.

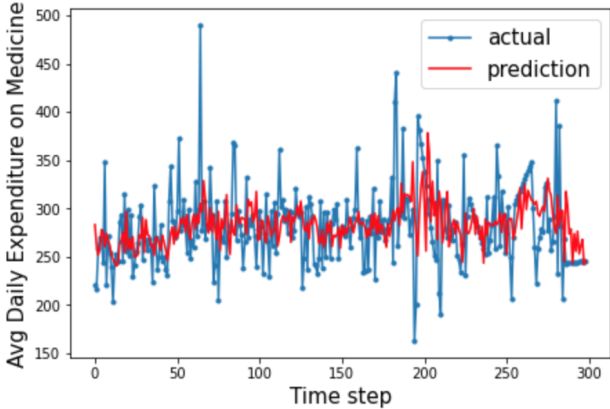
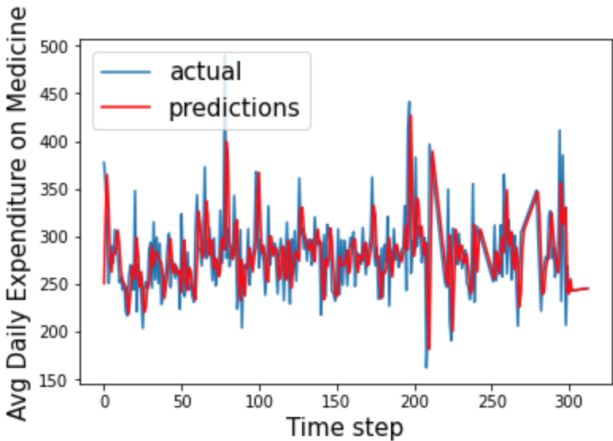
Problem 4.

A non-Deep Learning, ARIMA model was trained on the Medical Expenditures dataset for comparison purposes. ARIMA is an acronym for: AutoRegressive Integrated Moving Average. These models are specifically used to forecast a time series using the series past values and are characterized by three terms: lag order (p), degree of differencing (d), and order of the moving average (q). In the table below, these three factors were modified and the mean square error (MSE) was recorded for each.

p	d	q	Test MSE
1	0	0	2513
0	1	0	2565
0	0	1	4413
1	1	1	1608
1	1	0	2116
2	1	1	1605
3	1	1	1600
7	1	1	1613
3	1	2	1606
5	2	1	Matrix non-invertible
1	0	7	Matrix non-invertible
1	0	12	Matrix non-invertible
0	0	7	2379
2	0	1	Did not converge

It would be assumed that increasing the moving average (q) component would result in a lower MSE, but that is not indicated from the data collected above. Instead, increasing the autoregressive (p) component resulted in less error. But increasing this factor too much raised the error. This particular model also had difficulties with certain p, d, q combinations and would result in non-convergence or a 'matrix non-invertible' error. The best trial run based on lowest MSE was 3, 1, 1 for p, d, q, respectively.

In the following table the ARIMA model prediction and MSE is compared to the Deep Learning model prediction and RMSE.

	Deep Learning (DL) Model	ARIMA Model
RMSE/ MSE	Train RMSE: 23.317861557006836 Test RMSE: 32.44575119018555	Test MSE: 1600.211
Prediction		

Although the RMSE is much lower for the DL model, the ARIMA model better predicts the actual dataset as shown in the prediction graphs. The DL model has a smoother curve and does not account for the variance of the actual dataset. I would say the ARIMA model, for the reason of better prediction accuracy to the actual dataset, is a better model. However, both models do not have a very low MSE or RMSE and are still not “good enough”.